

# FYS4150 H20 Project 1

Olav Fønstelien

September 9, 2020

## Abstract

In this assignment we investigate and compare three different algorithms for a numerical solution to Poisson's equation with Dirichlet boundary conditions on a domain  $x \in \mathbb{R}$ . The efficiency of the three algorithms are compared with regards to speed and memory. We will see that the efficiency of a purely library-based general algorithm for solving a system of linear equations is greatly surpassed when we take into account the tridiagonal character of the coefficient matrix; and that the efficiency is further increased when we take into account the fixed pattern of the diagonals in the equation. We will also see that the relative error of the numerical solution has a simple analytical representation which only depends on step size  $h$ .

I have implemented the algorithms in C++ using the Armadillo library. See `project1.cpp`, the `Makefile` and a `README.md` at my repository <https://github.com/fonstelien/FYS4150/tree/master/project1>.

## A: Establishing the system of linear equations

Poisson's equation is a second degree differential equation on the form  $u''(x) = f(x, u(x))$ , and the assignment already gives us the form of the numerical approximation  $v_i \approx u_i$  to this problem as

$$-\frac{v_{i-1} - 2v_i + v_{i+1}}{h^2} = f_i \Rightarrow -v_{i-1} + 2v_i - v_{i+1} = h^2 f_i. \quad (1)$$

With boundary values  $u_0 = u(0) = \alpha$  and  $u_{n+1} = u(1) = \beta$ , we get the equations

$$\begin{aligned} -\alpha + 2v_1 - v_2 &= h^2 f_1 \\ -v_1 + 2v_2 - v_3 &= h^2 f_2 \\ &\vdots \\ -v_{n-1} + 2v_n - \beta &= h^2 f_n. \end{aligned}$$

Now, with  $\alpha = \beta = 0$  we can represent this as the set of linear equations  $\mathbf{A}\mathbf{v} = \tilde{\mathbf{b}}$  where  $\mathbf{A} \in \mathbb{R}^{n \times n}$  is tridiagonal with diagonal weights  $-1, 2$  and  $-1$ ,  $\mathbf{v} = \{v_1, v_2, \dots, v_n\}$ , and  $\tilde{\mathbf{b}} = \{h^2 f_1, h^2 f_2, \dots, h^2 f_n\}$ :

$$\begin{bmatrix} 2 & -1 & 0 & \cdots & 0 \\ -1 & 2 & -1 & \ddots & \vdots \\ 0 & \ddots & \ddots & \ddots & 0 \\ \vdots & \ddots & -1 & 2 & -1 \\ 0 & \cdots & 0 & -1 & 2 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} = h^2 \begin{bmatrix} f_1 \\ f_2 \\ \vdots \\ f_n \end{bmatrix}$$

## B: General numerical solution to a tridiagonal system of linear equations

The first programming task is to find an algorithm for a solution to the general tridiagonal problem

$$\mathbf{A}\mathbf{v} = \begin{bmatrix} b_1 & c_1 & 0 & \cdots & 0 \\ a_2 & b_2 & c_2 & \ddots & \vdots \\ 0 & \ddots & \ddots & \ddots & 0 \\ \vdots & \ddots & a_{n-1} & b_{n-1} & c_{n-1} \\ 0 & \cdots & 0 & a_n & b_n \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} = \tilde{\mathbf{b}} = \begin{bmatrix} \tilde{b}_1 \\ \tilde{b}_2 \\ \vdots \\ \tilde{b}_n \end{bmatrix}.$$

To solve this, let's first subtract  $\frac{a_2}{b_1}$  times the first row from the second row in  $\mathbf{A}$  to get rid of the first non-zero coefficient:

$$\left[ \begin{array}{ccccc|c} b_1 & c_1 & 0 & \cdots & 0 & \tilde{b}_1 \\ 0 & b_2 - \frac{a_2}{b_1}c_1 & c_2 & \ddots & \vdots & \tilde{b}_2 - \frac{a_2}{b_1}\tilde{b}_1 \\ \vdots & \ddots & \ddots & \ddots & \vdots & \vdots \end{array} \right].$$

We will denote these new values in the central diagonal and constant vector  $b_2^* = b_2 - \frac{a_2}{b_1}c_1$  and  $\tilde{b}_2^* = \tilde{b}_2 - \frac{a_2}{b_1}\tilde{b}_1$ , with the general terms

$$\begin{aligned} b_i^* &= b_i - \frac{a_i}{b_{i-1}}c_{i-1} \\ \tilde{b}_i^* &= \tilde{b}_i - \frac{a_i}{b_{i-1}}\tilde{b}_{i-1} \end{aligned}, \text{ for } i = 1, 2, \dots, n.$$

Forward substitution of the rest of the rows will then give us  $v_n$ :

$$\left[ \begin{array}{cccc|c} b_1 & c_1 & 0 & \cdots & 0 & \tilde{b}_1 \\ 0 & b_2^* & c_2 & \ddots & \vdots & \tilde{b}_2^* \\ \vdots & \ddots & b_3^* & \ddots & & \tilde{b}_3^* \\ & & & \ddots & & \vdots \\ & & & & b_{n-1}^* & c_{n-1} & \tilde{b}_{n-1}^* \\ & & & & 0 & b_n^* & \tilde{b}_n^* \end{array} \right] \Rightarrow v_n = \tilde{b}_n^*/b_n^* = \tilde{b}_n^{**}.$$

Now that we have found  $v_n$ , finding the rest is easy. First we find  $v_{n-1} = \tilde{b}_{n-1}^* - c_{n-1}\tilde{b}_n^{**} = \tilde{b}_{n-1}^{**}$ , and then the remaining  $v_i$  in the same manner:

$$\begin{bmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \ddots & \vdots \\ \vdots & \ddots & \ddots & 0 \\ 0 & \cdots & 0 & 1 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} = \begin{bmatrix} \tilde{b}_1^{**} \\ \tilde{b}_2^{**} \\ \vdots \\ \tilde{b}_n^{**} \end{bmatrix}$$

Counting the number of FLOPs needed to solve the tridiagonal  $\mathbf{A}\mathbf{v} = \tilde{\mathbf{b}}$  problem, we see that updating each term in the forward substitution process requires one division, one multiplication and one subtraction, and that we must do this  $n - 1$  times; that is  $N_{\text{fwd}} = 2 \times 3 \times (n - 1) = 6(n - 1)$ . Solving for  $x_n$  requires one operation, and backward substitution to solve for remaining  $n - 1$  requires 3 operations each, giving  $N_{\text{bwd}} = 1 + 3 \times (n - 1) \approx 3(n - 1)$  operations. All in all, the required number of FLOPs for a system of size  $n$  is

$$N_{\text{tridiag}} = N_{\text{fwd}} + N_{\text{bwd}} = 9(n - 1).$$

We should note that the above algorithm allows us to reuse the space which originally contained the description of the problem; that is  $\mathbf{A}$  and  $\tilde{\mathbf{b}}$ ; and also that we have found an algorithm which will solve a set of  $n$  equation in  $\mathcal{O}(n)$  time. This means that it is objectively efficient and scales well. My C++ implementation:

---

```
/* General tridiagonal system solver*/
colvec solve_tridiag(rowvec a, rowvec b, rowvec c, colvec b_tilde) {
    int n = b_tilde.n_rows;
    colvec solution = b_tilde;
    rowvec b_mod = b;

    // Forward substitution
    for (int i = 1; i < n; i++) {
        double fac = a[i]/b_mod[i-1];
        b_mod[i] -= fac*c[i-1];
        solution[i] -= fac*solution[i-1];
    }
}
```

```

// Solving for the last row
solution[n-1] /= b_mod[n-1];

// Backward substitution gives solution
for (int i = n-2; i > -1; i--)
    solution[i] = (solution[i] - c[i]*solution[i+1])/b_mod[i];

return solution;
}

```

---

Let's move on and compare how the numerical approximation given in equation 1 performs relative to the closed-form solution. The assignment gives us the functions  $f$  and the closed-form solution  $u$  to the differential equation:

$$f(x) = 100e^{-10x}$$

$$u(x) = 1 - (1 - e^{-10})x - e^{-10x}.$$

We insert  $f_i = f(x_i) = f(hi)$  into our numerical approximation for various step sizes  $h$  and observe from the plotted results in Figure 1 below that the approximation  $v_i$  approaches  $u_i$  for decreasing  $h$ , as expected.




Figure 1: Plot of the numerical solution to the Poisson equation vs the closed-form solution. When  $n > 100$  we see that the numerical solution is very close to the closed-form.

## C: Optimized numerical solution to the Poisson equation

Let's now take into consideration the fixed form of the diagonals of the coefficient matrix

$$\mathbf{A} = \begin{bmatrix} 2 & -1 & 0 & \cdots & 0 \\ -1 & 2 & -1 & \ddots & \vdots \\ 0 & \ddots & \ddots & \ddots & 0 \\ \vdots & \ddots & -1 & 2 & -1 \\ 0 & \cdots & 0 & -1 & 2 \end{bmatrix}.$$

In contrast to the general case solved in Section B, this time we know everything there is to know about  $\mathbf{A}$ , and the solution will only be dependent on the constant vector  $\tilde{\mathbf{b}}$ . As soon as we have found out which steps are needed to transform  $\mathbf{A}$  to  $\mathbf{I}$ , we can concentrate only on operations on  $\tilde{\mathbf{b}}$ , which will yield an algorithm which is optimized both with regards to speed and memory.

To find the solution, we first multiply row 2 with 2 and add row 1:

$$\left[ \begin{array}{ccccc|c} 2 & -1 & 0 & \cdots & 0 & \tilde{b}_1 \\ 0 & 3 & -2 & \ddots & \vdots & 2\tilde{b}_2 + \tilde{b}_1 \\ 0 & -1 & 2 & -1 & \vdots & \tilde{b}_3 \\ \vdots & \ddots & \ddots & \ddots & \vdots & \vdots \end{array} \right].$$

We now see that to get rid of  $a_3 = -1$  in the next row, we must this time multiply row 3 with 3 and add row 2, which will mean that in row 4 we must multiply with 4, and so on until we get  $[\mathbf{A}|\tilde{\mathbf{b}}]$  on the wanted form, and we can solve for  $v_n$ :

$$\left[ \begin{array}{cccccc|c} 2 & -1 & 0 & \cdots & & 0 & \tilde{b}_1 \\ 0 & 3 & -2 & \ddots & & \vdots & \tilde{b}_2^* \\ \vdots & \ddots & 4 & -3 & & & \tilde{b}_3^* \\ & & & \ddots & \ddots & \vdots & \vdots \\ \vdots & & & & \ddots & n & -(n-1) & \tilde{b}_{n-1}^* \\ 0 & \cdots & & & & 0 & n+1 & \tilde{b}_n^* \end{array} \right] \Rightarrow v_n = \tilde{b}_n^*/(n+1).$$

Here we have denoted the new values in the constant vector  $\tilde{\mathbf{b}}$  as  $\tilde{b}_2^* = 2\tilde{b}_2 - \tilde{b}_1$ ,  $\tilde{b}_3^* = 3\tilde{b}_3 - \tilde{b}_2^*$ , with the general term:

$$\tilde{b}_i^* = i\tilde{b}_i + \tilde{b}_{i-1}^*$$

We now find  $v_{n-1} = \frac{1}{n}(\tilde{b}_{n-1}^* + (n-1)\tilde{b}_n^*) = \tilde{b}_{n-1}^{**}$ , and again on general form  $v_i = \frac{1}{i+1}(\tilde{b}_i^* + i\tilde{b}_{i+1}^*) = \tilde{b}_i^{**}$ , which gives us the solution

$$\mathbf{v} = \tilde{\mathbf{b}}^{**}.$$

The resulting C++ implementation of this algorithm very much resembles `solve_tridiag()` from Section B, but it's a bit simpler since there is less stuff to be done in the optimized Poisson algorithm, as seen in `solve_poisson()` below:

---

```

/* Optimized Poisson system solver*/
colvec solve_poisson(colvec b_tilde) {
    int n = b_tilde.n_rows;
    colvec solution = b_tilde;

    // Forward substitution
    for (int i = 1; i < n; i++)
        solution[i] = (i+1)*solution[i] + solution[i-1];

    // Solving for the last row
    solution[n-1] /= (n+1);

    // Backward substitution gives solution
    for (int i = n-2; i > -1; i--)
        solution[i] = (solution[i] + (i+1)*solution[i+1])/(i+2);

    return solution;
}

```

---

As mentioned above, this algorithm is optimized both with regards to speed and memory. Memory required to solve a given problem is limited to the memory needed to describe it. Counting the FLOPs, we see that since we know how the forward substitution will unfold, each  $n-1$  rows require one multiplication and one addition, and  $N_{\text{fwd}} = 2(n-1)$ . Likewise for backward substitution, we only need one multiplication, one addition, and one division, and  $N_{\text{bwd}} = 3(n-1)$ . In total, the number of FLOPs required for a system of size  $n$  is

$$N_{\text{Poisson}} = N_{\text{fwd}} + N_{\text{bwd}} = 5(n-1).$$

Compared to the general case discussed in Section B, this of course is quite good, but the real benefit comes from the greatly decreased memory use. Numerical problems such as these are usually memory-bound, not CPU bound. We see that general triagonal system required three vectors  $\mathbf{a}, \mathbf{b}, \mathbf{c}$  for the coefficients and a fourth  $\tilde{\mathbf{b}}$  for the constants which must all be streamed from main memory, or at least Level 3 cache for a system of some size. The optimized Poisson system, however, only needs the constant vector  $\tilde{\mathbf{b}}$ , and therefore requires far fewer reads and writes for solving the system. This is the reason for the approximately 2.5-3 times speedup that we see in Table 1 below.

Table 1: CPU time needed to solve the Poisson equation. The algorithm optimized for the Poisson equation has better speedup than expected from counting the FLOPs. The reason is fewer read/write operations. Results are averaged over 5 runs.

n	General time [ms]	Optimized time [ms]	Speedup factor
100	0.103	0.042	2.452381
1000	0.659	0.277	2.379061
10000	10.179	3.385	3.007090
100000	55.765	20.122	2.771345
1000000	459.470	189.190	2.428617

## D: Relative error of the numerical solution to the Poisson equation

Given the closed-form solution  $u(x)$  to the Poisson equation  $-u''(x) = f(x)$ , we can calculate the relative error of our approximation  $v_i$  to  $u_i$ :

$$\varepsilon_i = \left| \frac{v_i - u_i}{u_i} \right|.$$

If we calculate this value for a sample of  $i$ , we will quickly notice that for a given step size  $h$ ,  $\varepsilon_i$  seems to be constant. To investigate this further, let's assume for a moment that indeed it *is* constant, such that  $\varepsilon_i = \varepsilon$ , and

$$\varepsilon = \left| \frac{v_{i-1} - u_{i-1}}{u_{i-1}} \right| = \left| \frac{v_i - u_i}{u_i} \right| = \left| \frac{v_{i+1} - u_{i+1}}{u_{i+1}} \right| \Rightarrow \begin{cases} v_{i-1} &= u_{i-1}(1 + \varepsilon) \\ v_i &= u_i(1 + \varepsilon) \\ v_{i+1} &= u_{i+1}(1 + \varepsilon) \end{cases}.$$

If we install this for  $v_{i-1}, v_i, v_{i+1}$  in the numerical approximation  $2v_i - v_{i-1} - v_{i+1} = h^2 f_i$ , we get that the relative error is given as

$$\varepsilon = \left| 1 + \frac{h^2 f_i}{2u_i - u_{i-1} - u_{i+1}} \right|.$$

Now, with  $f_i = f(ih) = 100e^{-10ih}$  and  $u_i = u(ih) = 1 - (1 - e^{-10})ih - e^{-10ih}$ , we get an expression for  $\varepsilon$  that is independent of  $i$ , meaning that the relative error is indeed constant for any step size  $h$ :

$$\begin{aligned} \varepsilon &= 1 + \frac{h^2 \cdot 100e^{-10ih}}{2e^{-10ih} - e^{-10(i-1)h} - e^{-10(i+1)h}} = 1 + \frac{100h^2}{2 - e^{10h} - e^{-10h}} \\ &= 1 - \frac{50h^2}{\cosh 10h - 1}. \end{aligned} \quad (2)$$

It is also interesting to see what happens as  $h \rightarrow 0$ .  $\cosh 0 = 1$  and  $\sinh 0 = 0$ , so by applying L'Hopital's twice we will see that the limit of the relative error is – luckily – zero:

$$\lim_{h \rightarrow 0} \varepsilon = 1 - \lim_{h \rightarrow 0} \frac{50h^2}{\cosh 10h - 1} = 1 - \lim_{h \rightarrow 0} \frac{10h}{\sinh 10h} = 1 - \lim_{h \rightarrow 0} \frac{1}{\cosh 10h} = 1 - 1 = 0.$$

From Table 2 we see that this mathematical error fits well with the errors that we find when we calculate it from the numerical solution. The difference is due to the round-off error stemming from the limited precision in a `double float`. We also see that up to about  $n = 10^6$ , the computer error decreases, before it starts to increase, meaning that we will reach maximum accuracy in this region of  $n$ . The C++ implementation is given further below.

Table 2: Relative error for decreasing step size  $h$  on a  $\log_{10}$  scale. The Computed error follows the Mathematical error from Equation 2 closely up until  $n = 10^5$ , after which it starts to deviate, and even increase for about  $n = 10^6$ . The mathematical error is calculated with the `decimal` module in Python.

n	Computed error	Mathematical error
10	-1.1797	-1.179698
100	-3.0880	-3.088037
1000	-5.0801	-5.080052
10000	-7.0793	-7.079268
100000	-9.0776	-9.079190
1000000	-10.1230	-11.079182
10000000	-9.0901	-13.080922
100000000	-8.1294	

---

```

/* Calculates the log10 relative error */
colvec log_eps(colvec exact, colvec approximated) {
    colvec eps;

    // Check if exact solution vector has zeros
    if (min(abs(exact)) == 0.) {
        cerr << "error: cannot calculate relative error; " \
              << "exact solution contains one or more zeros." << endl;
        return eps;
    }

    eps = abs((approximated - exact)/exact);
    return log10(eps);
}

```

---



## E: Numerical solution by the LU decomposition

If  $\det \mathbf{A} \neq 0$ , the decomposition of  $\mathbf{A}$  into a lower-triangular matrix  $\mathbf{L}$  and an upper-triangular matrix  $\mathbf{U}$ , where  $\mathbf{A}, \mathbf{L}, \mathbf{U} \in \mathbb{R}^{n \times n}$ , exists, such that

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} = \mathbf{L}\mathbf{U}, \text{ and} \quad (3)$$

$$\mathbf{L}\mathbf{U} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ l_{21} & 1 & 0 & 0 \\ l_{31} & l_{32} & 1 & 0 \\ l_{41} & l_{42} & l_{43} & 1 \end{bmatrix} \begin{bmatrix} u_{11} & u_{12} & u_{13} & u_{14} \\ 0 & u_{22} & u_{23} & u_{24} \\ 0 & 0 & u_{33} & u_{34} \\ 0 & 0 & 0 & u_{44} \end{bmatrix}.$$

A tridiagonal matrix with diagonals  $-1, 2, -1$  is obviously invertible, and hence  $\det \mathbf{A} \neq 0$ . We will not discuss how to decompose  $\mathbf{A}$  into  $\mathbf{L}\mathbf{U}$ , since that is not in the scope of the assignment, but concentrate on the computational complexity of solving the set of equations which emerges from Equation 3. In general we have that  $\mathbf{A}\mathbf{v} = \tilde{\mathbf{b}}$ , into which we insert  $\mathbf{A} = \mathbf{L}\mathbf{U}$ , such that

$$\mathbf{A}\mathbf{v} = \mathbf{L}\mathbf{U}\mathbf{v} = \mathbf{L}(\mathbf{U}\mathbf{v}) = \tilde{\mathbf{b}} \Rightarrow \begin{cases} \mathbf{U}\mathbf{v} = \mathbf{y} \\ \mathbf{L}\mathbf{y} = \tilde{\mathbf{b}} \end{cases}.$$

We begin by solving  $\mathbf{L}\mathbf{y} = \tilde{\mathbf{b}}$ . Let's first get rid of  $l_{21}$  in row 2, which requires one multiplication and one subtraction, so 2 FLOPs:

$$\left[ \begin{array}{cccc|c} 1 & 0 & 0 & 0 & \tilde{b}_1 \\ 0 & 1 & 0 & 0 & \tilde{b}_2 - l_{21}\tilde{b}_1 \\ l_{31} & l_{32} & 1 & 0 & \tilde{b}_3 \\ l_{41} & l_{42} & l_{43} & 1 & \tilde{b}_4 \end{array} \right]$$

Next, we get rid of  $l_{31}$  and  $l_{32}$  in row 3, which requires two multiplications and two subtractions, or 4 FLOPs:

$$\left[ \begin{array}{cccc|c} 1 & 0 & 0 & 0 & \tilde{b}_1 \\ 0 & 1 & 0 & 0 & \tilde{b}_2^* \\ 0 & 0 & 1 & 0 & \tilde{b}_3 - l_{31}\tilde{b}_1 - l_{32}\tilde{b}_2^* \\ l_{41} & l_{42} & l_{43} & 1 & \tilde{b}_4 \end{array} \right]$$

The next row requires 6 FLOPs, and so on. For a problem of size  $n$ , we need  $N_L = 2 + 4 + 6 + \dots + 2(n-1) = \frac{1}{2}(n-1) \cdot 2n = n(n-1)$  FLOPs for the "L" operation. Then, we need to solve the next set of equations, which we do in much the same manner, except for that it requires an extra division operation in each row:

$$\left[ \begin{array}{cccc|c} u_{11} & u_{12} & u_{13} & u_{14} & \tilde{b}_1^* \\ 0 & u_{22} & u_{23} & u_{24} & \tilde{b}_2^* \\ 0 & 0 & u_{33} & u_{34} & \tilde{b}_3^* \\ 0 & 0 & 0 & 1 & \tilde{b}_4^*/u_{44} \end{array} \right] \sim \left[ \begin{array}{cccc|c} u_{11} & u_{12} & u_{13} & u_{14} & \tilde{b}_1^* \\ 0 & u_{22} & u_{23} & u_{24} & \tilde{b}_2^* \\ 0 & 0 & 1 & 0 & (\tilde{b}_3^* - u_{34}\tilde{b}_4^{**})/u_{33} \\ 0 & 0 & 0 & 1 & \tilde{b}_4^{**} \end{array} \right]$$

Continuing in the same fashion we see that for the "U" operation we need  $N_U = 1 + 3 + 5 + \dots + (2n - 1) = \frac{1}{2}n \cdot 2n = n^2$  FLOPs. Adding all together, we get that the number of FLOPs needed for solving a set of linear equations with size  $n$  and *known* LU decomposition of  $\mathbf{A}$  is

$$N_{LU} = N_L + N_U = n(n - 1) + n^2 = n(2n - 1).$$

The C++ implementation of the algorithm with Armadillo is very simple, but to ensure that the library takes advantage of the diagonal shape of the matrices, we need to pass this information in to `arma::solve()` with the `arma::trimatl()` and `arma::trimatu()` arguments:

---

```
/* Solver for LU decomposed system without permutation matrix */
colvec solve_LU(mat L, mat U, colvec b_tilde) {
    colvec y;

    y = solve(trimatl(L), b_tilde);
    return solve(trimatu(U), y);
}
```

---

Here we should note that the solution to a problem of size  $n$  can be solved in  $\mathcal{O}(n^2)$  time by LU decomposition, which is a factor  $n$  better than the straightforward row reduction runs in  $\mathcal{O}(n^3)$  time. However, as Table 3 shows, compared to the optimized Poisson algorithm discussed in Section C, the LU method is still slow. Description of the problem also requires  $\mathcal{O}(n^2)$  space, in contrast to  $\mathcal{O}(n)$  for the specialized algorithm, and for problems of size  $n > 10^4$  or maybe  $10^5$ , the problem cannot at all be solved on a regular computer with the LU method due to memory restrictions.

Table 3: CPU time required for the solution by LU matrix decomposition algorithm relative to the optimized Poisson algorithm. Results are averaged over 5 runs.

n	LU	Poisson	Speedup
100	4.312	0.042	102.666667
1000	219.650	0.277	792.960289
10000	1649.400	3.385	487.267356

## Discussion and conclusion

This exercise illustrates very well that great improvements both with regards to speed and memory usage can be achieved with optimized algorithms. We also see that if the problem at hand is *large*, a simple solution with standardized library functions may not even be possible due to resource limitations.